

House Prices Analysis and Prediction

May 1, 2022

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.api import OLS

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms

%matplotlib inline
```

```
[2]: import warnings
warnings.filterwarnings("ignore")
```

```
[3]: BASE_COLOR = sns.color_palette()[0]
TEST_SIZE = 0.2
RANDOM_STATE = 42
BATCH_SIZE = 32
LR = 0.01
```

1 Load train and test Datasets

```
[4]: train_df = pd.read_csv("./data/train.csv")
test_df = pd.read_csv("./data/test.csv")
```

```
[5]: train_df.head()
```

```
[5]:   Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
0    1           60      RL          65.0     8450   Pave   NaN      Reg
1    2           20      RL          80.0     9600   Pave   NaN      Reg
2    3           60      RL          68.0    11250   Pave   NaN      IR1
3    4           70      RL          60.0     9550   Pave   NaN      IR1
4    5           60      RL          84.0    14260   Pave   NaN      IR1

      LandContour Utilities  ... PoolArea PoolQC Fence MiscFeature MiscVal MoSold  \
0             Lvl1   AllPub  ...         0    NaN   NaN           NaN         0      2
1             Lvl1   AllPub  ...         0    NaN   NaN           NaN         0      5
2             Lvl1   AllPub  ...         0    NaN   NaN           NaN         0      9
3             Lvl1   AllPub  ...         0    NaN   NaN           NaN         0      2
4             Lvl1   AllPub  ...         0    NaN   NaN           NaN         0     12

      YrSold  SaleType  SaleCondition  SalePrice
0     2008         WD         Normal     208500
1     2007         WD         Normal     181500
2     2008         WD         Normal     223500
3     2006         WD      Abnorml     140000
4     2008         WD         Normal     250000
```

[5 rows x 81 columns]

2 Wrangle train Dataset

2.1 Dropping NaN values

dropping NaN value with as follows: - if column has NaN values more than 200 entry will drop that column - if column has NaN values less than 200: - if column is numeric will fillna with the **mean** - if column is categorical will dropit

```
[6]: nan_cols = train_df.isna().sum()
nan_cols = nan_cols[nan_cols > 0]
nan_cols
```

```
[6]: LotFrontage      259
Alley              1369
MasVnrType         8
MasVnrArea         8
BsmtQual          37
BsmtCond          37
BsmtExposure      38
BsmtFinType1      37
BsmtFinType2      38
Electrical         1
```

```

FireplaceQu      690
GarageType       81
GarageYrBlt      81
GarageFinish     81
GarageQual       81
GarageCond       81
PoolQC          1453
Fence           1179
MiscFeature     1406
dtype: int64

```

```

[7]: nan_cols_to_drop = nan_cols[nan_cols > 200]
train_df = train_df.drop(nan_cols_to_drop.index.values, axis='columns')
train_df.head()

```

```

[7]:   Id  MSSubClass MSZoning  LotArea Street LotShape LandContour Utilities \
0    1           60       RL   8450   Pave       Reg         Lvl     AllPub
1    2           20       RL   9600   Pave       Reg         Lvl     AllPub
2    3           60       RL  11250   Pave      IR1         Lvl     AllPub
3    4           70       RL   9550   Pave      IR1         Lvl     AllPub
4    5           60       RL  14260   Pave      IR1         Lvl     AllPub

```

```

      LotConfig LandSlope  ... EnclosedPorch 3SsnPorch ScreenPorch PoolArea \
0     Inside      Gtl  ...              0          0          0          0
1      FR2      Gtl  ...              0          0          0          0
2     Inside      Gtl  ...              0          0          0          0
3     Corner      Gtl  ...             272          0          0          0
4      FR2      Gtl  ...              0          0          0          0

```

```

      MiscVal  MoSold  YrSold  SaleType  SaleCondition  SalePrice
0          0        2   2008        WD        Normal    208500
1          0        5   2007        WD        Normal    181500
2          0        9   2008        WD        Normal    223500
3          0        2   2006        WD      Abnorml    140000
4          0       12   2008        WD        Normal    250000

```

```
[5 rows x 75 columns]
```

```

[8]: nan_cols = train_df.isna().sum()
nan_cols = nan_cols[nan_cols > 0]
# nan_cols
non_numeric_nan_cols = train_df[nan_cols.index.values].select_dtypes('object').
    ↪columns
non_numeric_nan_cols

```

```

[8]: Index(['MasVnrType', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1',
        'BsmtFinType2', 'Electrical', 'GarageType', 'GarageFinish',

```

```
'GarageQual', 'GarageCond'],
dtype='object')
```

```
[9]: train_df = train_df.drop(non_numeric_nan_cols, axis="columns")
train_df.head()
```

```
[9]:
```

	Id	MSSubClass	MSZoning	LotArea	Street	LotShape	LandContour	Utilities	\
0	1	60	RL	8450	Pave	Reg	Lvl	AllPub	
1	2	20	RL	9600	Pave	Reg	Lvl	AllPub	
2	3	60	RL	11250	Pave	IR1	Lvl	AllPub	
3	4	70	RL	9550	Pave	IR1	Lvl	AllPub	
4	5	60	RL	14260	Pave	IR1	Lvl	AllPub	

	LotConfig	LandSlope	...	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea	\
0	Inside	Gtl	...	0	0	0	0	
1	FR2	Gtl	...	0	0	0	0	
2	Inside	Gtl	...	0	0	0	0	
3	Corner	Gtl	...	272	0	0	0	
4	FR2	Gtl	...	0	0	0	0	

	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	0	2	2008	WD	Normal	208500
1	0	5	2007	WD	Normal	181500
2	0	9	2008	WD	Normal	223500
3	0	2	2006	WD	Abnorml	140000
4	0	12	2008	WD	Normal	250000

[5 rows x 64 columns]

2.2 Checking any duplicates

```
[10]: train_df[train_df.duplicated()]
```

```
[10]: Empty DataFrame
Columns: [Id, MSSubClass, MSZoning, LotArea, Street, LotShape, LandContour,
Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType,
HouseStyle, OverallQual, OverallCond, YearBuilt, YearRemodAdd, RoofStyle,
RoofMatl, Exterior1st, Exterior2nd, MasVnrArea, ExterQual, ExterCond,
Foundation, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, Heating, HeatingQC,
CentralAir, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath,
BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, KitchenQual,
TotRmsAbvGrd, Functional, Fireplaces, GarageYrBlt, GarageCars, GarageArea,
PavedDrive, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch,
PoolArea, MiscVal, MoSold, YrSold, SaleType, SaleCondition, SalePrice]
Index: []
```

[0 rows x 64 columns]

No Duplicated row in the dataset

```
[11]: nan_cols = train_df.isna().sum()
      nan_cols = nan_cols[nan_cols > 0]
      nan_cols
```

```
[11]: MasVnrArea      8
      GarageYrBlt    81
      dtype: int64
```

```
[12]: train_df['MasVnrArea'].fillna(train_df['MasVnrArea'].mean(), inplace=True)
      train_df['GarageYrBlt'].fillna(train_df['GarageYrBlt'].mean(), inplace=True)

      nan_cols = train_df.isna().sum()
      nan_cols = nan_cols[nan_cols > 0]

      assert len(nan_cols) == 0
```

2.3 Inspecting Categorical Columns

```
[13]: numeric_cols = train_df.select_dtypes('number').columns
      categorical_cols = train_df.select_dtypes('object').columns
```

```
[14]: numeric_cols
```

```
[14]: Index(['Id', 'MSSubClass', 'LotArea', 'OverallQual', 'OverallCond',
            'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',
            'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
            'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
            'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
            'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
            'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
            'MoSold', 'YrSold', 'SalePrice'],
            dtype='object')
```

```
[15]: categorical_cols
```

```
[15]: Index(['MSZoning', 'Street', 'LotShape', 'LandContour', 'Utilities',
            'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
            'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
            'Exterior2nd', 'ExterQual', 'ExterCond', 'Foundation', 'Heating',
            'HeatingQC', 'CentralAir', 'KitchenQual', 'Functional', 'PavedDrive',
            'SaleType', 'SaleCondition'],
            dtype='object')
```

2.3.1 MSZoning Analysis

```
[16]: train_df['MSZoning'].value_counts(normalize=True).sort_values(ascending=False)
```

```
[16]: RL          0.788356
      RM          0.149315
      FV          0.044521
      RH          0.010959
      C (all)     0.006849
      Name: MSZoning, dtype: float64
```

MSZoning has a high bias towards the **Residential Low Density** and **Residential Medium Density** zoning types.

2.3.2 HouseStyle Analysis

```
[17]: train_df['HouseStyle'].value_counts(normalize=True).sort_values(ascending=False)
```

```
[17]: 1Story      0.497260
      2Story      0.304795
      1.5Fin      0.105479
      SLvl       0.044521
      SFoyer      0.025342
      1.5Unf      0.009589
      2.5Unf      0.007534
      2.5Fin      0.005479
      Name: HouseStyle, dtype: float64
```

HouseStyle has major proportions towards **One Story**, **Two Story** and **One and one-half story**: **2nd level unfinished**

2.3.3 Functional Analysis

```
[18]: train_df['Functional'].value_counts(normalize=True).sort_values(ascending=False)
```

```
[18]: Typ          0.931507
      Min2       0.023288
      Min1       0.021233
      Mod        0.010274
      Maj1       0.009589
      Maj2       0.003425
      Sev        0.000685
      Name: Functional, dtype: float64
```

Functional column has large proportion towards **Typical Functionality**

2.3.4 SaleCondition Analysis

```
[19]: train_df['SaleCondition'].value_counts(normalize=True).  
      ↪sort_values(ascending=False)
```

```
[19]: Normal      0.820548  
      Partial    0.085616  
      Abnorml     0.069178  
      Family     0.013699  
      Alloca     0.008219  
      AdjLand    0.002740  
      Name: SaleCondition, dtype: float64
```

SaleCondition has large proportion in being **Normal**

2.3.5 Heating Analysis

```
[20]: train_df['Heating'].value_counts(normalize=True).sort_values(ascending=False)
```

```
[20]: GasA      0.978082  
      GasW     0.012329  
      Grav     0.004795  
      Wall     0.002740  
      OthW     0.001370  
      Floor    0.000685  
      Name: Heating, dtype: float64
```

Heating columns has large proportion towards **Gas forced warm air furnace**

2.4 Cleaning Categorical Columns

```
[21]: ## MSZoning Column  
      train_df['zone_low_density'] = train_df['MSZoning'] == "RL"  
  
      ## HouseStyle Column  
      train_df['one_story_style'] = train_df['HouseStyle'] == '1Story'  
      train_df['two_story_style'] = train_df['HouseStyle'] == '2Story'  
      train_df['one_half_story_style'] = train_df['HouseStyle'] == '1.5Fin'  
  
      ## Functional Column  
      train_df['typical_functionality'] = train_df['Functional'] == 'Typ'  
  
      ## Sale Condition  
      train_df['sale_normal_condition'] = train_df['SaleCondition'] == 'Normal'  
  
      ## Heating Column  
      train_df['gas_heating'] = train_df['Heating'] == 'GasA'
```

```

## conver to numeric values
cols = ['zone_low_density', 'one_story_style', 'two_story_style',
        ↪ 'typical_functionality',
        ↪ 'one_half_story_style', 'gas_heating', 'sale_normal_condition']

train_df[cols] = train_df[cols].astype('int')

train_df[
    ['zone_low_density', 'one_story_style', 'two_story_style',
    ↪ 'typical_functionality',
    ↪ 'one_half_story_style', 'gas_heating', 'sale_normal_condition']
].head()

```

```

[21]:
zone_low_density  one_story_style  two_story_style  typical_functionality \
0                1                0                1                1
1                1                1                0                1
2                1                0                1                1
3                1                0                1                1
4                1                0                1                1

one_half_story_style  gas_heating  sale_normal_condition
0                    0            1                    1
1                    0            1                    1
2                    0            1                    1
3                    0            1                    0
4                    0            1                    1

```

2.5 Inspecting Numerical Columns

```
[22]: numeric_cols
```

```

[22]: Index(['Id', 'MSSubClass', 'LotArea', 'OverallQual', 'OverallCond',
            'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',
            'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
            'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
            'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
            'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
            'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
            'MoSold', 'YrSold', 'SalePrice'],
           dtype='object')

```

2.5.1 MSSubClass Analysis


- 20 1-STORY 1946 & NEWER ALL STYLES
- 30 1-STORY 1945 & OLDER
- 40 1-STORY W/FINISHED ATTIC ALL AGES
- 45 1-1/2 STORY - UNFINISHED ALL AGES

- 50 1-1/2 STORY FINISHED ALL AGES
- **60 2-STORY 1946 & NEWER**
- 70 2-STORY 1945 & OLDER
- 75 2-1/2 STORY ALL AGES
- 80 SPLIT OR MULTI-LEVEL
- 85 SPLIT FOYER
- 90 DUPLEX - ALL STYLES AND AGES
- 120 1-STORY PUD (Planned Unit Development) - 1946 & NEWER
- 150 1-1/2 STORY PUD - ALL AGES
- 160 2-STORY PUD - 1946 & NEWER
- 180 PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
- 190 2 FAMILY CONVERSION - ALL STYLES AND AGES

this column must be converted to be **categorical column**

```
[23]: train_df['MSSubClass'] = train_df['MSSubClass'].astype('object')


train_df['MSSubClass'].value_counts(normalize=True).plot(kind='bar')
plt.title("Count of each House Class")
plt.xlabel('Class')
plt.ylabel('Count');
```



2.5.2 YearBuilt Analysis

```
[24]: (train_df['YearBuilt'].value_counts()  
      [train_df['YearBuilt'].value_counts() > 20]  
      ).plot(kind='bar')
```

[24]: <AxesSubplot:>



2.5.3 YearRemodAdd Analysis

```
[25]: train_df['time_taken_to_remodel'] = train_df['YearRemodAdd'] -  
      ↪ train_df['YearBuilt']  
      train_df['time_taken_to_remodel'].describe()
```

```
[25]: count    1460.000000  
      mean      13.597945  
      std       24.476465  
      min        0.000000  
      25%        0.000000  
      50%        0.000000  
      75%       20.000000  
      max      123.000000  
      Name: time_taken_to_remodel, dtype: float64
```

```
[26]: train_df['time_taken_to_sell'] = train_df['YrSold'] - train_df['YearBuilt']
train_df['time_taken_to_sell'].describe()
```

```
[26]: count      1460.000000
mean         36.547945
std          30.250152
min           0.000000
25%           8.000000
50%          35.000000
75%          54.000000
max          136.000000
Name: time_taken_to_sell, dtype: float64
```

```
[27]: train_df['time_taken_to_sell_after_remodel'] = train_df['YrSold'] -
↳ train_df['YearRemodAdd']
train_df['time_taken_to_sell_after_remodel'].describe()
```

```
[27]: count      1460.000000
mean         22.950000
std          20.640653
min          -1.000000
25%           4.000000
50%          14.000000
75%          41.000000
max           60.000000
Name: time_taken_to_sell_after_remodel, dtype: float64
```

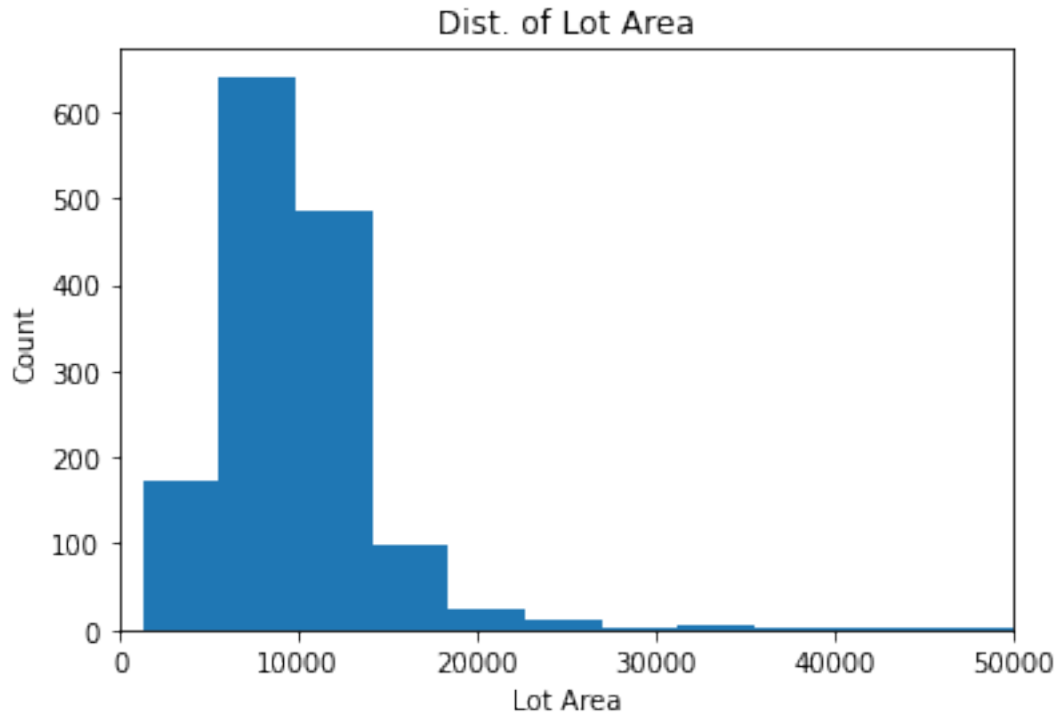
on average it took **36 Years** to sell a house and **22 Year** to sell a house that has been **ReModeled**

2.5.4 LotArea Analysis

```
[28]: train_df['LotArea'].plot(kind='hist', bins=50)
plt.title("Dist. of Lot Area")
plt.xlabel("Lot Area")
plt.ylabel('Count')
plt.xlim([0, 50_000]);

print(train_df['LotArea'].describe())
```

```
count      1460.000000
mean      10516.828082
std       9981.264932
min       1300.000000
25%       7553.500000
50%       9478.500000
75%      11601.500000
max      215245.000000
Name: LotArea, dtype: float64
```



most of the houses has a Lot Area around **5,000** to **15,000** square feet

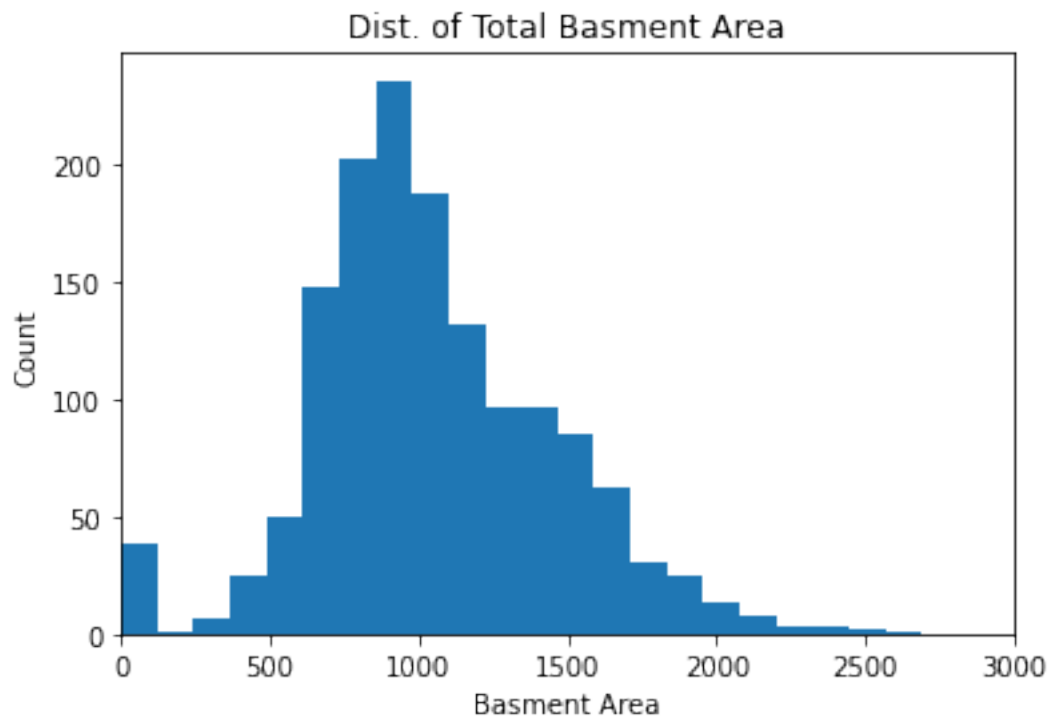
2.5.5 TotalBsmtSF

```
[29]: print(train_df['TotalBsmtSF'].describe())

train_df['TotalBsmtSF'].plot(kind='hist', bins=50)
plt.title("Dist. of Total Basement Area")
plt.xlabel("Basement Area")
plt.ylabel('Count');

plt.xlim([0, 3000]);
```

```
count    1460.000000
mean      1057.429452
std        438.705324
min         0.000000
25%        795.750000
50%        991.500000
75%       1298.250000
max       6110.000000
Name: TotalBsmtSF, dtype: float64
```



most of the houses has a Basment Area from **500** to **1500**

2.5.6 Bedroom and Bathroom Analysis

```
[30]: train_df['BedroomAbvGr'].value_counts()
```

```
[30]: 3    804
      2    358
      4    213
      1     50
      5     21
      6      7
      0      6
      8      1
      Name: BedroomAbvGr, dtype: int64
```


```
[31]: train_df['FullBath'].value_counts()
```

```
[31]: 2    768
      1    650
      3     33
      0      9
      Name: FullBath, dtype: int64
```

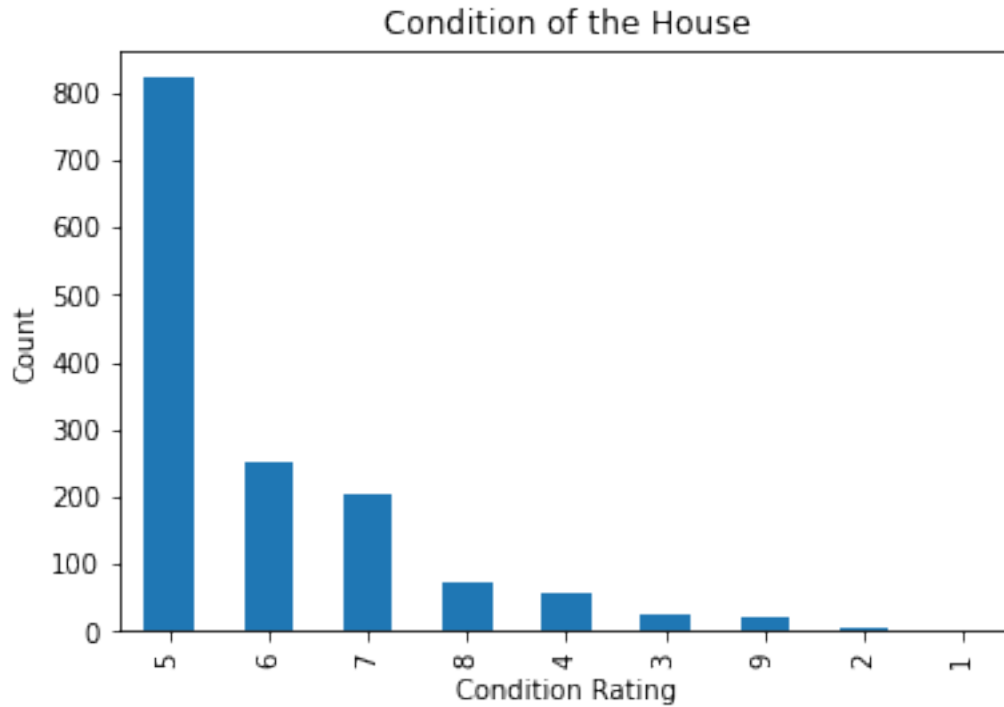
most of the house have a **3, 2 or 4** bedrooms above the grade and **2 or 1** bathroom above the grade

2.5.7 OverallQual and OverallCond Analysis

```
[32]: train_df['OverallQual'].value_counts().plot(kind='bar');  
plt.title("Quality of the House")  
plt.xlabel("Quality Rating")  
plt.ylabel("Count");
```



```
[33]: train_df['OverallCond'].value_counts().plot(kind='bar')  
plt.title("Condition of the House")  
plt.xlabel("Condition Rating")  
plt.ylabel("Count");
```




most of the houses have a rating condition and quality around **5 or 6**

2.5.8 SalePrice Analysis

```
[34]: print(train_df['SalePrice'].describe())

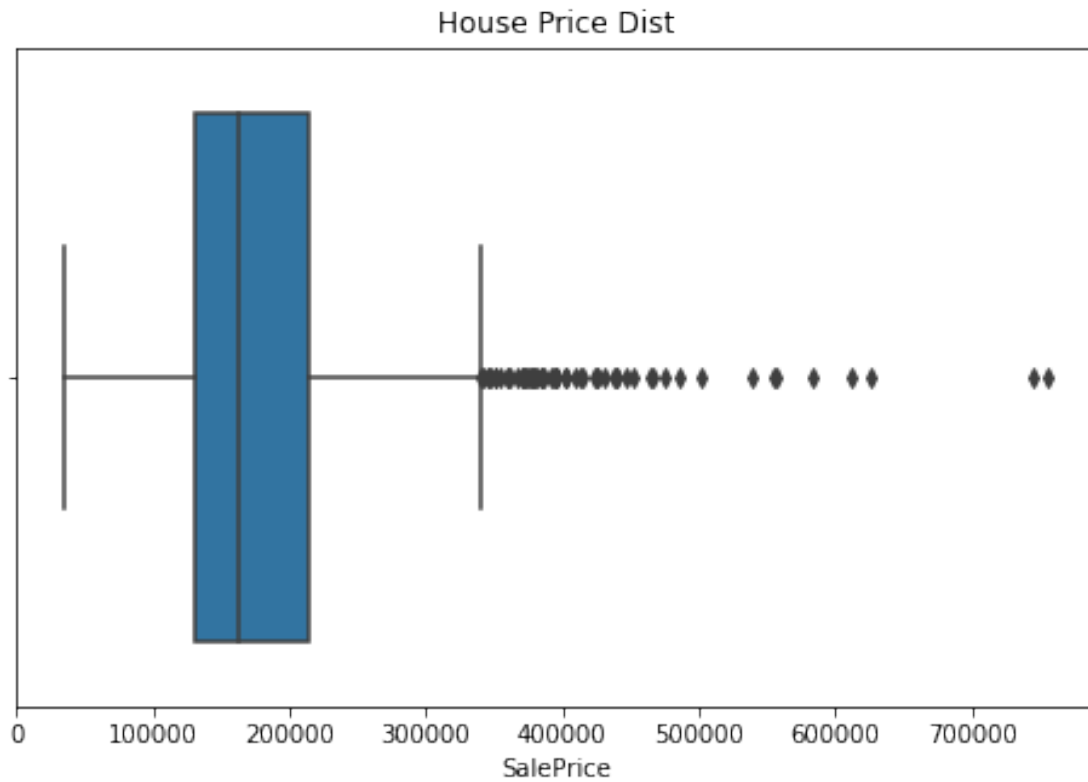
train_df['SalePrice'].plot(kind='hist', bins=30)
plt.title("Price of House")
plt.xlabel("Price")
plt.ylabel("Count");
```

```
count      1460.000000
mean       180921.195890
std        79442.502883
min         34900.000000
25%        129975.000000
50%        163000.000000
75%        214000.000000
max         755000.000000
Name: SalePrice, dtype: float64
```



most houses has price ranging between **1,000,000** to **2,500,000** us dollars


```
[35]: plt.figure(figsize=(8, 5))
sns.boxplot(data=train_df, x='SalePrice')
plt.title("House Price Dist");
```

2.5.9 Correlation Matrix in the Dataset

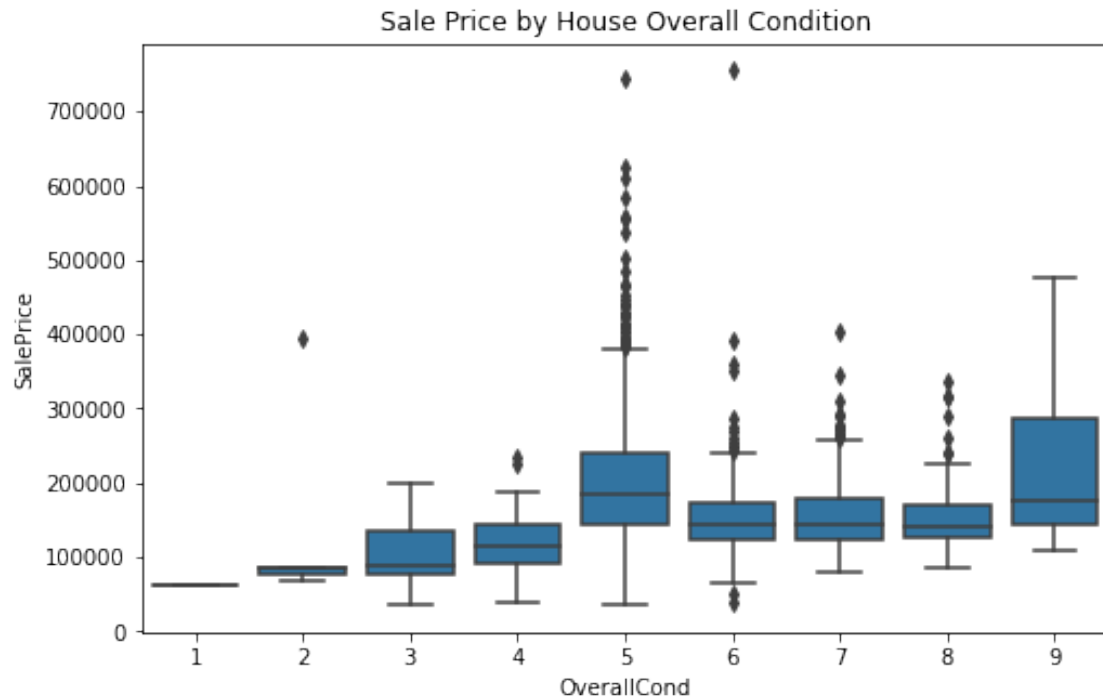
```
[36]: corr = train_df.corr()

fig, ax = plt.subplots(nrows=1, figsize=(15, 14))
sns.heatmap(corr, ax=ax)
ax.set_title("Correlation Matrix");
```




2.5.10 SalePrice vs OverallCond

```
[37]: plt.figure(figsize=(8, 5))
sns.boxplot(data=train_df, x='OverallCond', y='SalePrice', color=BASE_COLOR);
plt.title("Sale Price by House Overall Condition");
```



2.5.11 SalePrice vs TotalBsmtSF

```
[38]: plt.figure(figsize=(8, 5))
sns.regplot(data=train_df, x='TotalBsmtSF', y='SalePrice', color=BASE_COLOR);
plt.title("Sale Price by House Basement Area");
```



the relation shows existence of outlier point at **6000** so, it would be better if we remove it

```
[39]: train_df.drop(train_df[train_df['TotalBsmtSF'] > 5000].index, inplace=True)
      train_df[train_df['TotalBsmtSF'] > 5000]
```


[39]: Empty DataFrame

Columns: [Id, MSSubClass, MSZoning, LotArea, Street, LotShape, LandContour, Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType, HouseStyle, OverallQual, OverallCond, YearBuilt, YearRemodAdd, RoofStyle, RoofMatl, Exterior1st, Exterior2nd, MasVnrArea, ExterQual, ExterCond, Foundation, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, Heating, HeatingQC, CentralAir, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, KitchenQual, TotRmsAbvGrd, Functional, Fireplaces, GarageYrBlt, GarageCars, GarageArea, PavedDrive, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, MiscVal, MoSold, YrSold, SaleType, SaleCondition, SalePrice, zone_low_density, one_story_style, two_story_style, one_half_story_style, typical_functionality, sale_normal_condition, gas_heating, time_taken_to_remodel, time_taken_to_sell, time_taken_to_sell_after_remodel]

Index: []


[0 rows x 74 columns]

```
[40]: plt.figure(figsize=(8, 5))
sns.regplot(data=train_df, x='TotalBsmtSF', y='SalePrice',
            line_kws=dict(color='r'));
plt.title("Sale Price by House Basement Area");
```



2.5.12 SalePrice vs GrLivArea

```
[41]: plt.figure(figsize=(8, 5))
sns.regplot(data=train_df, x='GrLivArea', y='SalePrice',
            line_kws=dict(color='r'));
plt.title("Sale Price by House Basement Area");
```




there is outliers in the GrLivArea greater than 4000 so removing it

```
[42]: train_df.drop(train_df[train_df['GrLivArea'] > 4000].index, inplace=True)
      train_df[train_df['GrLivArea'] > 4000]
```

```
[42]: Empty DataFrame
Columns: [Id, MSSubClass, MSZoning, LotArea, Street, LotShape, LandContour,
Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType,
HouseStyle, OverallQual, OverallCond, YearBuilt, YearRemodAdd, RoofStyle,
RoofMatl, Exterior1st, Exterior2nd, MasVnrArea, ExterQual, ExterCond,
Foundation, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, Heating, HeatingQC,
CentralAir, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath,
BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, KitchenQual,
TotRmsAbvGrd, Functional, Fireplaces, GarageYrBlt, GarageCars, GarageArea,
PavedDrive, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch,
PoolArea, MiscVal, MoSold, YrSold, SaleType, SaleCondition, SalePrice,
zone_low_density, one_story_style, two_story_style, one_half_story_style,
typical_functionality, sale_normal_condition, gas_heating,
time_taken_to_remodel, time_taken_to_sell, time_taken_to_sell_after_remodel]
Index: []


[0 rows x 74 columns]
```

```
[43]: plt.figure(figsize=(8, 5))
sns.regplot(data=train_df, x='GrLivArea', y='SalePrice',
            line_kws=dict(color='r'));
plt.title("Sale Price by House Ground Area");
```



2.5.13 GarageArea vs SalePrice

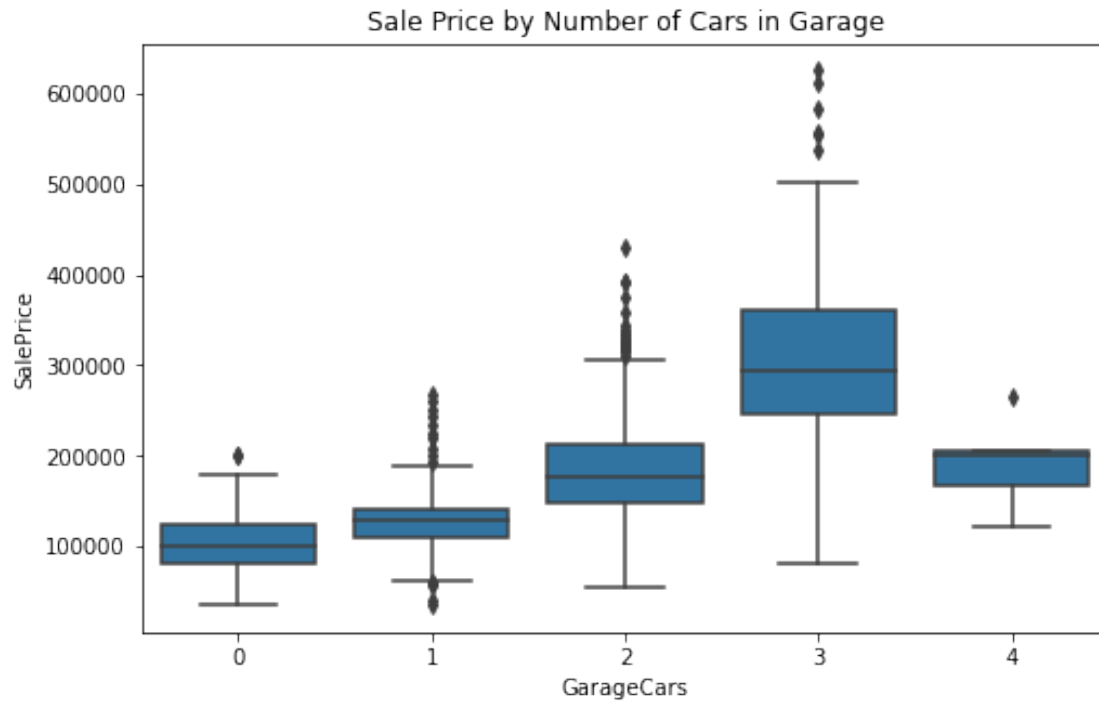
```
[44]: plt.figure(figsize=(8, 5))
sns.regplot(data=train_df, x='GarageArea', y='SalePrice',
            line_kws=dict(color='r'));
plt.title("Sale Price by Garage Area");
```



positive correlation indicating an increase in price as the garage area increases.


2.5.14 GarageCars vs SalePrice

```
[45]: plt.figure(figsize=(8, 5))
sns.boxplot(data=train_df, x='GarageCars', y='SalePrice', color=BASE_COLOR);
plt.title("Sale Price by Number of Cars in Garage");
```


positive correlation indicating an increase in price as number of cars garage can fit increases

```
[46]: ## GarageCars vs GarageArea  
## False correlation or colinearity in the dataset  
## Does make sense  
  
plt.figure(figsize=(8, 5))  
sns.boxplot(data=train_df, x='GarageCars', y='GarageArea', color=BASE_COLOR);
```



2.5.15 SalePrice vs FullBath

```
[47]: plt.figure(figsize=(8, 5))
sns.boxplot(data=train_df, x='FullBath', y='SalePrice', color=BASE_COLOR);
plt.title("Sale Price by Number of Bathrooms");
```



```
[ ]:
```

3 Statmodels OLS Model Analysis

```
[48]: ols_df = train_df.select_dtypes('number')
      ols_df.head()
```

```
[48]:
```

	Id	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	\
0	1	8450	7	5	2003	2003	196.0	
1	2	9600	6	8	1976	1976	0.0	
2	3	11250	7	5	2001	2002	162.0	
3	4	9550	7	5	1915	1970	0.0	
4	5	14260	8	5	2000	2000	350.0	

	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF	...	zone_low_density	one_story_style	\
0	706	0	150	...	1	0	
1	978	0	284	...	1	1	
2	486	0	434	...	1	0	
3	216	0	540	...	1	0	
4	655	0	490	...	1	0	

	two_story_style	one_half_story_style	typical_functionality	\
0	1	0	1	

1	0	0	1
2	1	0	1
3	1	0	1
4	1	0	1

	sale_normal_condition	gas_heating	time_taken_to_remodel \
0	1	1	0
1	1	1	0
2	1	1	1
3	0	1	55
4	1	1	0

	time_taken_to_sell	time_taken_to_sell_after_remodel
0	5	5
1	31	31
2	7	6
3	91	36
4	8	8

[5 rows x 46 columns]

```
[49]: ols_df = ols_df.drop(['YrSold', 'YearBuilt', 'YearRemodAdd', 'Id'],
    ↪axis='columns')
ols_df
```

```
[49]:
```

	LotArea	OverallQual	OverallCond	MasVnrArea	BsmtFinSF1	BsmtFinSF2 \
0	8450	7	5	196.0	706	0
1	9600	6	8	0.0	978	0
2	11250	7	5	162.0	486	0
3	9550	7	5	0.0	216	0
4	14260	8	5	350.0	655	0
...
1455	7917	6	5	0.0	0	0
1456	13175	6	6	119.0	790	163
1457	9042	7	9	0.0	275	0
1458	9717	5	6	0.0	49	1029
1459	9937	5	6	0.0	830	290

	BsmtUnfSF	TotalBsmtSF	1stFlrSF	2ndFlrSF	...	zone_low_density \
0	150	856	856	854	...	1
1	284	1262	1262	0	...	1
2	434	920	920	866	...	1
3	540	756	961	756	...	1
4	490	1145	1145	1053	...	1
...
1455	953	953	953	694	...	1
1456	589	1542	2073	0	...	1

1457	877	1152	1188	1152 ...	1
1458	0	1078	1078	0 ...	1
1459	136	1256	1256	0 ...	1

	one_story_style	two_story_style	one_half_story_style	\
0	0	1	0	
1	1	0	0	
2	0	1	0	
3	0	1	0	
4	0	1	0	
...	
1455	0	1	0	
1456	1	0	0	
1457	0	1	0	
1458	1	0	0	
1459	1	0	0	

	typical_functionality	sale_normal_condition	gas_heating	\
0	1	1	1	
1	1	1	1	
2	1	1	1	
3	1	0	1	
4	1	1	1	
...	
1455	1	1	1	
1456	0	1	1	
1457	1	1	1	
1458	1	1	1	
1459	1	1	1	

	time_taken_to_remodel	time_taken_to_sell	\
0	0	5	
1	0	31	
2	1	7	
3	55	91	
4	0	8	
...	
1455	1	8	
1456	10	32	
1457	65	69	
1458	46	60	
1459	0	43	

	time_taken_to_sell_after_remodel
0	5
1	31
2	6

```

3          36
4          8
...
1455       7
1456      22
1457       4
1458      14
1459      43

```

[1456 rows x 42 columns]

```
[50]: ols_df['intercept'] = 1
```

```
[51]: ols_model = OLS(ols_df['SalePrice'], ols_df.drop('SalePrice', axis='columns'))
result = ols_model.fit()
result.summary()
```

```
[51]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  SalePrice    R-squared:                0.870
Model:                            OLS      Adj. R-squared:            0.866
Method:                 Least Squares    F-statistic:                248.8
Date:                Sun, 01 May 2022    Prob (F-statistic):          0.00
Time:                  01:35:18    Log-Likelihood:            -16959.
No. Observations:                  1456    AIC:                      3.400e+04
Df Residuals:                      1417    BIC:                      3.420e+04
Df Model:                            38
Covariance Type:                  nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
LotArea                0.5892    0.082      7.168    0.000
0.428    0.750
OverallQual            1.4e+04    980.392    14.279    0.000
1.21e+04    1.59e+04
OverallCond            4810.6610    864.986     5.562    0.000
3113.871    6507.451
MasVnrArea             28.7187     4.871     5.896    0.000
19.163    38.274
BsmtFinSF1             20.2892     2.144     9.464    0.000
16.084    24.495
BsmtFinSF2             0.3289     3.650     0.090    0.928

```

-6.830	7.488				
BsmtUnfSF		-0.2817	1.969	-0.143	0.886
-4.145	3.582				
TotalBsmtSF		20.3364	2.826	7.196	0.000
14.793	25.880				
1stFlrSF		23.5409	5.308	4.435	0.000
13.129	33.953				
2ndFlrSF		28.9580	5.521	5.245	0.000
18.128	39.788				
LowQualFinSF		-18.3571	12.279	-1.495	0.135
-42.445	5.730				
GrLivArea		34.1418	4.674	7.305	0.000
24.974	43.310				
BsmtFullBath		1188.0415	2142.964	0.554	0.579
-3015.682	5391.765				
BsmtHalfBath		-6472.7349	3345.912	-1.935	0.053
-1.3e+04	90.739				
FullBath		-1839.3572	2303.181	-0.799	0.425
-6357.367	2678.653				
HalfBath		-1745.3871	2214.303	-0.788	0.431
-6089.052	2598.278				
BedroomAbvGr		-1.179e+04	1379.653	-8.545	0.000
-1.45e+04	-9082.100				
KitchenAbvGr		-2.403e+04	3983.327	-6.034	0.000
-3.18e+04	-1.62e+04				
TotRmsAbvGrd		5593.5981	1013.354	5.520	0.000
3605.763	7581.433				
Fireplaces		1842.6362	1457.101	1.265	0.206
-1015.671	4700.943				
GarageYrBlt		106.9927	56.099	1.907	0.057
-3.053	217.038				
GarageCars		3023.0111	2356.422	1.283	0.200
-1599.439	7645.461				
GarageArea		18.5384	8.052	2.302	0.021
2.743	34.334				
WoodDeckSF		15.2935	6.518	2.346	0.019
2.508	28.079				
OpenPorchSF		27.3756	12.513	2.188	0.029
2.829	51.922				
EnclosedPorch		11.2541	13.682	0.823	0.411
-15.585	38.093				
3SsnPorch		10.4082	25.354	0.411	0.681
-39.328	60.144				
ScreenPorch		38.0293	13.973	2.722	0.007
10.620	65.439				
PoolArea		-3.5168	21.287	-0.165	0.869
-45.275	38.241				

MiscVal	-1.0412	1.505	-0.692	0.489
-3.993	1.911			
MoSold	-16.9945	276.902	-0.061	0.951
-560.176	526.188			
zone_low_density	6488.6639	2135.912	3.038	0.002
2298.775	1.07e+04			
one_story_style	6570.5700	3002.646	2.188	0.029
680.461	1.25e+04			
two_story_style	-2732.2378	4015.387	-0.680	0.496
-1.06e+04	5144.504			
one_half_story_style	659.9872	3689.555	0.179	0.858
-6577.590	7897.564			
typical_functionality	1.371e+04	3210.058	4.272	0.000
7415.270	2e+04			
sale_normal_condition	-7827.3892	2033.137	-3.850	0.000
-1.18e+04	-3839.107			
gas_heating	-4403.7303	5333.956	-0.826	0.409
-1.49e+04	6059.570			
time_taken_to_remodel	-18.4268	29.553	-0.624	0.533
-76.399	39.546			
time_taken_to_sell	-226.4515	38.136	-5.938	0.000
-301.261	-151.642			
time_taken_to_sell_after_remodel	-208.0247	37.895	-5.490	0.000
-282.361	-133.689			
intercept	-2.603e+05	1.13e+05	-2.308	0.021
-4.81e+05	-3.91e+04			
=====				
Omnibus:	422.947	Durbin-Watson:	2.035	
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4012.154	
Skew:	1.071	Prob(JB):	0.00	
Kurtosis:	10.845	Cond. No.	1.01e+16	
=====				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3e-21. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

""

4 Sklearn Models Analysis

```
[52]: features_cols = [
        "time_taken_to_sell",
        "OverallQual",
        "OverallCond",
```



```

    "MasVnrArea",
    "TotalBsmtSF",
    "GrLivArea",
    "BedroomAbvGr",
    "TotRmsAbvGrd",
    "GarageArea",
    "ScreenPorch",
    "zone_low_density",
    "typical_functionality",
    "sale_normal_condition",
]

```

```

[53]: features_df = train_df[features_cols]
      features_df.head()

```

```

[53]:   time_taken_to_sell  OverallQual  OverallCond  MasVnrArea  TotalBsmtSF  \
0                5                7                5        196.0         856
1               31                6                8          0.0        1262
2                7                7                5        162.0         920
3               91                7                5          0.0         756
4                8                8                5        350.0        1145

      GrLivArea  BedroomAbvGr  TotRmsAbvGrd  GarageArea  ScreenPorch  \
0         1710              3              8         548           0
1         1262              3              6         460           0
2         1786              3              6         608           0
3         1717              3              7         642           0
4         2198              4              9         836           0

      zone_low_density  typical_functionality  sale_normal_condition
0                1                1                1
1                1                1                1
2                1                1                1
3                1                1                0
4                1                1                1

```

```

[54]: X = features_df.values
      y = train_df['SalePrice']

```

```

[55]: sc = StandardScaler()

      X = sc.fit_transform(X)
      X

```

```

[55]: array([[ -1.04612544,  0.66564559, -0.51746143, ...,  0.51903606,
              0.27156272,  0.46516048],
            [ -0.18625648, -0.06470858,  2.17654515, ...,  0.51903606,

```

```

    0.27156272, 0.46516048],
[-0.97998167, 0.66564559, -0.51746143, ..., 0.51903606,
 0.27156272, 0.46516048],
...,
[ 1.07047507, 0.66564559, 3.07454735, ..., 0.51903606,
 0.27156272, 0.46516048],
[ 0.77282812, -0.79506274, 0.38054077, ..., 0.51903606,
 0.27156272, 0.46516048],
[ 0.21060611, -0.79506274, 0.38054077, ..., 0.51903606,
 0.27156272, 0.46516048]])

```

```

[56]: x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=TEST_SIZE,
                                                    random_state=RANDOM_STATE)

assert x_train.shape[1] == x_test.shape[1]

```

4.1 Linear Model

```

[57]: lm = LinearRegression()
lm.fit(x_train, y_train)
lm_preds = lm.predict(x_test)

print("R2 Score for Linear Model : ", lm.score(x_test, y_test))
print("RMSE for Linear Model : ", np.sqrt(mean_squared_error(y_test, lm_preds)))

```

R2 Score for Linear Model : 0.8464733772982462

RMSE for Linear Model : 28386.122329701033

4.2 RandomForest Model

```

[58]: rf_model = RandomForestRegressor()
rf_model.fit(x_train, y_train)
rf_preds = rf_model.predict(x_test)

print("R2 Score for Random Forest Model : ", rf_model.score(x_test, y_test))
print("RMSE for Random Forest Model : ", np.sqrt(mean_squared_error(y_test,
↪rf_preds)))

```

R2 Score for Random Forest Model : 0.8718567738601649

RMSE for Random Forest Model : 25933.554655943077

4.3 DecisionTreeRegressor Model

```

[59]: dt_model = DecisionTreeRegressor()
dt_model.fit(x_train, y_train)
dt_preds = dt_model.predict(x_test)

```

```
print("R2 Score for Decision Tree Model : ", rf_model.score(x_test, y_test))
print("RMSE for Decision Tree Model : ", np.sqrt(mean_squared_error(y_test,
↳ dt_preds)))
```

R2 Score for Decision Tree Model : 0.8718567738601649

RMSE for Decision Tree Model : 39583.2371666652

Random Forest Model is the best till this moment with **RMSE = 25126.752** and **R2 = 0.8797** and I will use to predict the test dataframe

4.4 Generating Prediction on Test DataFrame

4.4.1 Cleaning Test Dataset like in Train Dataset

```
[60]: ## MSZoning Column
test_df['zone_low_density'] = test_df['MSZoning'] == "RL"

## HouseStyle Column
test_df['one_story_style'] = test_df['HouseStyle'] == '1Story'
test_df['two_story_style'] = test_df['HouseStyle'] == '2Story'
test_df['one_half_story_style'] = test_df['HouseStyle'] == '1.5Fin'

## Functional Column
test_df['typical_functionality'] = test_df['Functional'] == 'Typ'

## Sale Condition
test_df['sale_normal_condition'] = test_df['SaleCondition'] == 'Normal'

## Heating Column
test_df['gas_heating'] = test_df['Heating'] == 'GasA'

## conver to numeric values
cols = ['zone_low_density', 'one_story_style', 'two_story_style',
↳ 'typical_functionality',
      'one_half_story_style', 'gas_heating', 'sale_normal_condition']

test_df[cols] = test_df[cols].astype('int')

# test_df[
#     ['zone_low_density', 'one_story_style', 'two_story_style',
↳ 'typical_functionality',
#     'one_half_story_style', 'gas_heating', 'sale_normal_condition']
# ].head()
```

```
[61]: test_df['time_taken_to_sell'] = test_df['YrSold'] - test_df['YearBuilt']
```

```
[62]: test_features_df = test_df[features_cols]
test_features_df.head()
```

```
[62]:
```

	time_taken_to_sell	OverallQual	OverallCond	MasVnrArea	TotalBsmtSF	\
0	49	5	6	0.0	882.0	
1	52	6	6	108.0	1329.0	
2	13	5	5	0.0	928.0	
3	12	6	6	20.0	926.0	
4	18	8	5	0.0	1280.0	

	GrLivArea	BedroomAbvGr	TotRmsAbvGrd	GarageArea	ScreenPorch	\
0	896	2	5	730.0	120	
1	1329	3	6	312.0	0	
2	1629	3	6	482.0	0	
3	1604	3	7	470.0	0	
4	1280	2	5	506.0	144	

	zone_low_density	typical_functionality	sale_normal_condition
0	0	1	1
1	1	1	1
2	1	1	1
3	1	1	1
4	1	1	1

```
[63]: test_features_df.isna().sum()
```

```
[63]: time_taken_to_sell      0
OverallQual                  0
OverallCond                  0
MasVnrArea                   15
TotalBsmtSF                   1
GrLivArea                    0
BedroomAbvGr                 0
TotRmsAbvGrd                 0
GarageArea                    1
ScreenPorch                   0
zone_low_density              0
typical_functionality         0
sale_normal_condition         0
dtype: int64
```

```
[64]: test_features_df['MasVnrArea'].fillna(test_features_df['MasVnrArea'].mean(),
      ↪inplace=True)
test_features_df['TotalBsmtSF'].fillna(test_features_df['TotalBsmtSF'].mean(),
      ↪inplace=True)
test_features_df['GarageArea'].fillna(test_features_df['GarageArea'].mean(),
      ↪inplace=True)

assert len(test_features_df.isna().sum()[test_features_df.isna().sum() != 0])
      ↪== 0
```

```
[65]: X_test = test_features_df.values
X_test = sc.fit_transform(X_test)
X_test
```

```
[65]: array([[ 0.41378443, -0.75110125,  0.40076604, ..., -1.79693781,
           0.27416383,  0.46021084],
          [ 0.51239787, -0.05487716,  0.40076604, ...,  0.55650229,
           0.27416383,  0.46021084],
          [-0.76957685, -0.75110125, -0.49741776, ...,  0.55650229,
           0.27416383,  0.46021084],
          ...,
          [ 0.31517099, -0.75110125,  1.29894985, ...,  0.55650229,
           0.27416383, -2.17291708],
          [-0.73670571, -0.75110125, -0.49741776, ...,  0.55650229,
           0.27416383,  0.46021084],
          [-0.76957685,  0.64134693, -0.49741776, ...,  0.55650229,
           0.27416383,  0.46021084]])
```

4.4.2 Evaluating Best Machine Learning Model.

```
[66]: test_preds = rf_model.predict(X_test)
test_preds
```

```
[66]: array([129459.   , 154134.5 , 172650.33, ..., 151392.5 , 113305.   ,
           233236.3 ])
```

```
[67]: test_df['SalePrice'] = test_preds
test_df[['Id', 'SalePrice']].head()
```

```
[67]:      Id  SalePrice
0   1461   129459.00
1   1462   154134.50
2   1463   172650.33
3   1464   190920.00
4   1465   210673.91
```

```
[68]: ## exporting the predictions to csv
test_df[['Id', 'SalePrice']].to_csv('predictions_1.csv', index=False)
```

4.5 PyTorch Neural Network Model

```
[69]: class HouseDataset(Dataset):
        def __init__(self, x_data, y_data, transform=None):
            """
            Args:
                x_data (np.ndarray) : x_train data array
```

```

        y_data (np.ndarray) : y_train data array
        transform (callable, optional): Optional transform to be applied
            on a sample.
    """

    self.x_data = x_data
    self.y_data = y_data
    self.transform = transform

    def __len__(self):
        return self.x_data.shape[0]

    def __getitem__(self, idx):
        if torch.is_tensor(idx):
            idx = idx.tolist()

        if not isinstance(self.y_data, np.ndarray):
            self.y_data = np.array(self.y_data)

        x_sample = self.x_data[idx]
        y_sample = self.y_data[idx]

        if self.transform is not None:
            x_sample = self.transform(x_sample)

        return x_sample, y_sample

train_ds = HouseDataset(x_train, y_train)
train_dl = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True)

sample = next(iter(train_dl))
sample[0], sample[1]

```

```

[69]: (tensor([[ 1.4673e+00,  6.6565e-01,  1.2785e+00, -5.7750e-01, -7.8313e-01,
                2.7081e-01,  1.6591e-01, -3.1416e-01,  3.5123e-01, -2.7061e-01,
                -1.9266e+00,  2.7156e-01,  4.6516e-01],
               [-1.0131e+00,  2.8567e+00, -5.1746e-01,  3.4103e+00,  3.8405e+00,
                2.2682e+00, -1.0603e+00,  9.2712e-01,  1.5687e+00, -2.7061e-01,
                5.1904e-01,  2.7156e-01,  4.6516e-01],
               [ 5.0825e-01, -7.9506e-01, -5.1746e-01, -5.7750e-01, -2.5501e+00,
                -7.7017e-01,  1.6591e-01, -9.3479e-01, -2.2253e+00, -2.7061e-01,
                5.1904e-01, -3.6824e+00,  4.6516e-01],
               [ 7.8319e-02, -7.9506e-01,  3.8054e-01, -1.4761e-01, -1.1810e-01,
                -1.0158e+00,  1.6591e-01, -9.3479e-01,  2.0312e+00, -2.7061e-01,
                5.1904e-01,  2.7156e-01,  4.6516e-01],
               [ 1.4673e+00, -7.9506e-01,  1.2785e+00, -5.7750e-01, -3.0256e-01,

```

-3.8358e-01, 1.6591e-01, -3.1416e-01, 5.8662e-02, -2.7061e-01,
 -1.9266e+00, 2.7156e-01, 4.6516e-01],
 [2.7675e-01, -7.9506e-01, -5.1746e-01, -5.7750e-01, 2.1684e-01,
 3.8156e-01, 1.6591e-01, 9.2712e-01, -5.7839e-01, -2.7061e-01,
 5.1904e-01, -3.6824e+00, 4.6516e-01],
 [1.8973e+00, -6.4709e-02, 3.8054e-01, -5.7750e-01, -7.4915e-01,
 -4.5308e-02, 1.6591e-01, 1.5478e+00, -2.2253e+00, -2.7061e-01,
 -1.9266e+00, 2.7156e-01, 4.6516e-01],
 [1.0043e+00, -7.9506e-01, 1.2785e+00, -5.7750e-01, -6.2536e-01,
 -7.8225e-01, 1.6591e-01, -9.3479e-01, -2.9054e-01, 4.5852e+00,
 -1.9266e+00, 2.7156e-01, 4.6516e-01],
 [-7.1541e-01, 1.3960e+00, -5.1746e-01, -5.7750e-01, 6.9498e-01,
 -3.4129e-01, -1.0603e+00, -9.3479e-01, -4.5154e-02, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.2115e+00, 1.3960e+00, -5.1746e-01, 1.1194e+00, 1.2896e+00,
 1.3178e+00, 1.6591e-01, 9.2712e-01, 1.4979e+00, -2.7061e-01,
 5.1904e-01, 2.7156e-01, -2.1498e+00],
 [5.7440e-01, -7.9506e-01, -5.1746e-01, -5.7750e-01, 8.0663e-01,
 -2.4867e-01, 1.6591e-01, -3.1416e-01, -5.5480e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.8626e-01, -6.4709e-02, 1.2785e+00, 4.0672e-01, 2.2412e-01,
 -7.9232e-01, 1.6591e-01, -3.1416e-01, 1.5304e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.0461e+00, 6.6565e-01, -5.1746e-01, 1.0742e+00, -1.2781e-01,
 1.0863e+00, 1.6591e-01, 9.2712e-01, 1.8849e+00, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.1453e+00, 1.3960e+00, -5.1746e-01, -5.7750e-01, -8.3409e-01,
 -1.8625e-01, 1.6591e-01, -3.1416e-01, -3.2357e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.1453e+00, 1.3960e+00, -5.1746e-01, -5.7750e-01, 1.0081e+00,
 -8.1551e-02, 1.6591e-01, 3.0648e-01, 6.5324e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.1453e+00, -6.4709e-02, -5.1746e-01, 8.7054e-01, 7.5566e-01,
 -2.9096e-01, -1.0603e+00, -3.1416e-01, -5.4591e-02, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.1784e+00, -6.4709e-02, -5.1746e-01, -1.5893e-01, -3.8023e-01,
 8.5876e-01, 1.3921e+00, 1.5478e+00, 9.2694e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [2.1061e-01, -6.4709e-02, 3.8054e-01, -5.7750e-01, 2.0713e-01,
 -7.4601e-01, 1.6591e-01, -9.3479e-01, -4.1323e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [-1.5318e-01, -6.4709e-02, 3.8054e-01, 2.1440e-01, -8.7536e-01,
 -1.6209e-01, 1.6591e-01, -3.1416e-01, -1.4897e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],
 [1.7319e+00, -7.9506e-01, 1.2785e+00, -5.7750e-01, -5.6954e-01,
 -1.8223e-01, 1.6591e-01, 3.0648e-01, 4.9280e-01, -2.7061e-01,
 5.1904e-01, 2.7156e-01, 4.6516e-01],

```

[-1.0131e+00, -6.4709e-02, -5.1746e-01,  3.8409e-01, -4.9187e-01,
 -1.3259e+00, -2.2865e+00, -1.5554e+00, -2.4335e-01, -2.7061e-01,
 -1.9266e+00,  2.7156e-01,  4.6516e-01],
[-1.2115e+00,  6.6565e-01,  3.8054e-01, -5.7750e-01, -4.2877e-01,
  5.1243e-01,  1.6591e-01,  3.0648e-01,  5.0224e-01, -2.7061e-01,
 -1.9266e+00,  2.7156e-01, -2.1498e+00],
[-1.2115e+00, -6.4709e-02, -5.1746e-01, -1.9286e-01,  5.1295e-01,
 -4.9231e-01, -1.0603e+00, -9.3479e-01,  4.7393e-01, -2.7061e-01,
 -1.9266e+00,  2.7156e-01, -2.1498e+00],
[ 5.0825e-01, -7.9506e-01, -5.1746e-01, -5.7750e-01, -1.6180e+00,
 -2.9498e-01,  1.6591e-01, -3.1416e-01,  4.7393e-01, -2.7061e-01,
  5.1904e-01, -3.6824e+00,  4.6516e-01],
[-1.2115e+00,  1.3960e+00, -5.1746e-01, -1.7024e-01,  6.0518e-01,
 -3.8761e-01,  1.6591e-01, -3.1416e-01,  3.7955e-01, -2.7061e-01,
 -1.9266e+00,  2.7156e-01, -2.1498e+00],
[ 6.0747e-01, -7.9506e-01, -1.4155e+00, -5.7750e-01, -5.8168e-01,
 -2.4062e-01,  1.6591e-01,  3.0648e-01, -1.0172e+00, -2.7061e-01,
 -1.9266e+00,  2.7156e-01,  4.6516e-01],
[ 1.3681e+00, -7.9506e-01,  1.2785e+00, -5.7750e-01, -8.1468e-01,
 -1.2715e+00, -1.0603e+00, -9.3479e-01, -1.3759e+00, -2.7061e-01,
 -1.9266e+00,  2.7156e-01,  4.6516e-01],
[-1.0792e+00,  6.6565e-01, -5.1746e-01,  1.9226e+00, -4.4333e-01,
 -7.2588e-01,  1.6591e-01, -3.1416e-01, -1.5841e-01, -2.7061e-01,
  5.1904e-01,  2.7156e-01,  4.6516e-01],
[-1.1784e+00,  6.6565e-01, -5.1746e-01,  8.9868e-03,  9.1585e-01,
 -1.5806e-01,  1.6591e-01, -3.1416e-01,  3.9787e-02, -2.7061e-01,
 -1.9266e+00,  2.7156e-01, -2.1498e+00],
[ 2.4368e-01, -6.4709e-02,  3.8054e-01, -5.7750e-01,  1.1319e+00,
  7.6815e-01, -1.0603e+00, -3.1416e-01,  5.0224e-01, -2.7061e-01,
  5.1904e-01, -3.6824e+00,  4.6516e-01],
[-1.2011e-01, -6.4709e-02,  3.8054e-01,  8.6489e-01, -2.3443e-02,
 -7.6816e-01,  1.6591e-01, -3.1416e-01, -5.6423e-01, -2.7061e-01,
  5.1904e-01,  2.7156e-01, -2.1498e+00],
[-7.8155e-01, -6.4709e-02, -5.1746e-01,  1.6424e-02, -2.9770e-01,
  4.3189e-01,  1.6591e-01,  3.0648e-01, -2.6836e-03, -2.7061e-01,
  5.1904e-01,  2.7156e-01,  4.6516e-01]], dtype=torch.float64),
tensor([178000, 466500,  93000, 148500, 130000, 143900, 128000, 119900, 180000,
        325624, 165000, 148000, 284000, 178000, 201000, 194700, 214000, 140000,
        159500, 112000, 140000, 227875, 185000, 130000, 229456,  93000, 105000,
        187500, 182000, 190000, 125000, 201000]))

```

```

[70]: class NNet(nn.Module):
        def __init__(self, input_size: int, hidden_sizes: list, output_size: int=1)␣
        ↪-> None:
            super().__init__()

            self.fc1 = nn.Linear(input_size, hidden_sizes[0])

```



```

        self.fc2 = nn.Linear(hidden_sizes[0], hidden_sizes[1])
        self.fc3 = nn.Linear(hidden_sizes[1], hidden_sizes[2])
        self.fc4 = nn.Linear(hidden_sizes[2], output_size)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        out = self.fc4(x)

        return out

class RMSELoss(nn.Module):
    def __init__(self):
        super().__init__()

        self.mse = nn.MSELoss()

    def forward(self, y_pred, y_true):
        return torch.sqrt(self.mse(y_pred, y_true))

```

```

[71]: def train(train_dl, model, optimizer, criterion, epochs=100) -> list:
        losses = []

        print("{:<8}|{:>15}".format("Epoch", "Loss"))
        print("="*24)
        epoch_loss = 0

        for epoch in range(epochs):
            epoch_losses = []
            for x_batch, y_batch in train_dl:

                y_preds = model(x_batch.float())
                loss = criterion(y_preds, y_batch.float())

                optimizer.zero_grad()
                loss.backward()
                optimizer.step()

                epoch_loss += loss.item()
                epoch_losses.append(loss.item())

            losses.append(np.mean(epoch_losses))

            if epoch % 10 == 0:
                print("{:<8}|{:>15}".format(epoch, round(np.mean(epoch_losses), 3)))

```

```
return losses
```

```
[72]: ## initializing mode
n_feaures = x_train.shape[1]
hidden_sizes = [32, 16, 8]
output_size = 1

model = NNet(n_feaures, hidden_sizes, output_size)

print(model)

## training Model
optimizer = optim.Adam(model.parameters(), lr=LR)
criterion = RMSELoss()

losses = train(train_dl, model, optimizer, criterion, epochs=200)
```

```
NNet(
  (fc1): Linear(in_features=13, out_features=32, bias=True)
  (fc2): Linear(in_features=32, out_features=16, bias=True)
  (fc3): Linear(in_features=16, out_features=8, bias=True)
  (fc4): Linear(in_features=8, out_features=1, bias=True)
)
```

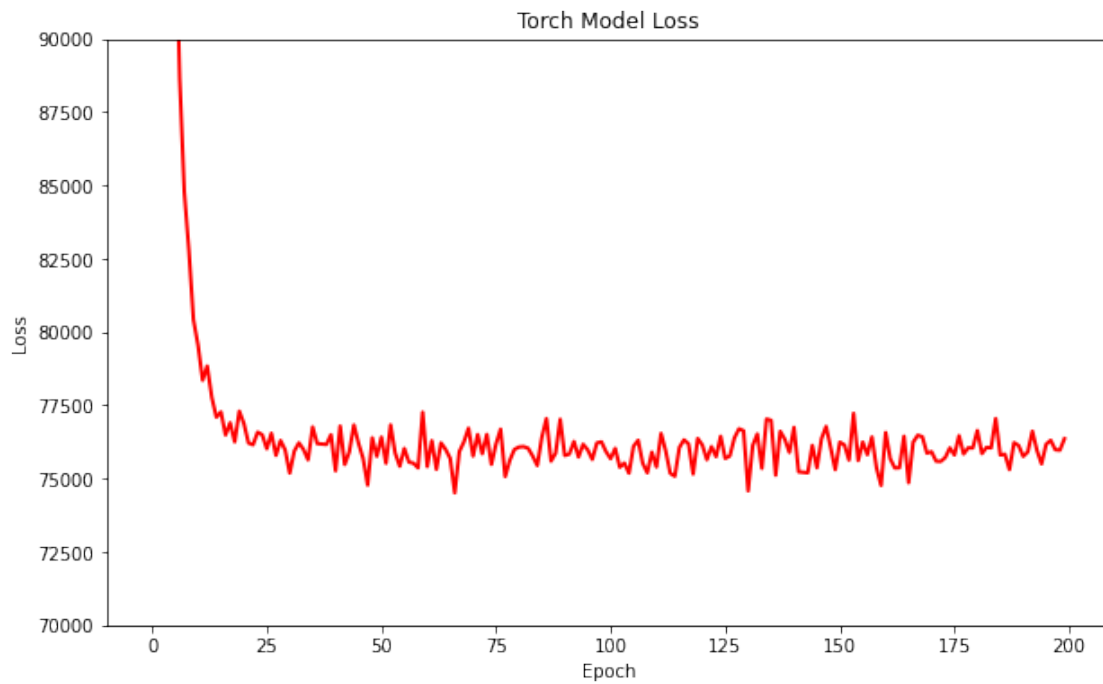
Epoch	Loss
0	196045.524
10	79553.088
20	76881.525
30	75187.745
40	75267.642
50	76417.62
60	75412.843
70	75768.498
80	76077.083
90	75788.585
100	75680.965
110	75396.359
120	76147.532
130	74582.772
140	76748.031
150	76243.72
160	76572.471
170	75917.212
180	76640.626
190	75758.835

```
[73]: %matplotlib inline

fig, ax = plt.subplots(nrows=1, figsize=(10, 6))

ax.plot(losses, 'r-', lw=2)
ax.set_title("Torch Model Loss")
ax.set_xlabel("Epoch")
ax.set_ylabel("Loss");
ax.set_ylim([70000, 90000])
```

[73]: (70000.0, 90000.0)



[]: